IoT-Based Intrusion Detection System Using Machine Learning

A Project Report submitted in partial fulfilment of the requirements for the degree of

Bachelor of Technology in Computer Engineering

Submitted by

Name	**MIS Number**
Mahi Patel	612303107
Malhar Falke	612303108
Yugandhar Pise	612303140
Vaibhav Tayade	612303184

Under the guidance of

Prof. S K Gaikwad

Department of Computer Engineering

COEP Technological University

Academic Year: 2025–26

Plagiarism Certificate

(Scanned copy of plagiarism percentage report signed by Project Guide to be inserted here)

DEPARTMENT OF COMPUTER ENGINEERING,

COEP TECHNOLOGICAL UNIVERSITY, PUNE – 411005

CERTIFICATE

Certified that the Sem VI project titled, "**IoT-Based Intrusion Detection System Using Machine Learning**" has been successfully completed by:

Mahi Patel	**612303107**
Malhar Falke	**612303108**
Yugandhar Pise	**612303140**
Vaibhav Tayade	**612303184**

and is approved for the partial fulfilment of the requirements for the degree of "B.Tech. Computer Engineering".

SIGNATURE

Prof. S K Gaikwad
Project Guide
Department of Computer Engineering
COEP Technological University,
Shivajinagar, Pune – 411005.

SIGNATURE

Prof. Pradeep Deshmukh
Head of Department
Department of Computer Engineering
COEP Technological University,
Shivajinagar, Pune – 411005.

Abstract

This project presents a lightweight, multi-dataset Intrusion Detection System (IoT-IDS) for IoT edge environments. The central challenge addressed is the cross-dataset generalisation gap: an IDS model that achieves near-perfect performance on the benchmark where it is trained can catastrophically fail when evaluated on traffic from a different network environment. In preliminary experiments, a model trained on CICIDS2017 and Bot-IoT achieved 99% and 98% accuracy respectively on their own test splits, but collapsed to 36.09% accuracy (F1 = 0.00) when evaluated on the unseen UNSW-NB15 dataset.

To resolve this, a Semantic Feature Alignment Framework (SFAF) is designed that maps heterogeneous feature schemas from three benchmark datasets — CICIDS2017, UNSW-NB15, and TON-IoT — to a common 12-dimensional representation. A unified XGBoost classifier trained on this aligned space achieves 98.33% accuracy (F1 = 0.9584, AUC = 0.9985) on CICIDS2017, 92.68% accuracy (F1 = 0.9432, AUC = 0.9844) on UNSW-NB15, and 99.41% accuracy (F1 = 0.9961, AUC = 0.9997) on TON-IoT, reducing the UNSW-NB15 generalisation gap by more than 56 percentage points. The unified model is compressed to a 45 KB ONNX artefact with sub-millisecond inference latency on Apple Silicon and a projected latency of 0.11 ms on Raspberry Pi 4, making it suitable for real-time IoT gateway deployment.

Contents

Abstract **ii**

List of Figures **vi**

List of Tables **vii**

1 Introduction **1**

1.1 Background 1
1.2 Motivation 1
1.3 Research Context 2
1.4 Objectives 2
1.5 Scope of the Project 2

2 Literature Review **4**

2.1 Intrusion Detection Systems 4
2.2 IoT-Specific Security Challenges 4
2.3 Machine Learning Approaches for IDS 4
2.4 Benchmark Datasets for IoT IDS 5
2.5 Cross-Dataset Generalisation Gap 5
2.6 Lightweight and Edge-Deployable IDS 5
2.7 Summary of Literature Gaps 6

3 Research Gaps and Problem Statement **7**

3.1 Identified Research Gaps 7
3.2 Problem Statement 7
3.3 Project Objectives 8

4 Proposed Methodology and Solution **9**

4.1 System Overview 9
4.2 Dataset Preprocessing 9
4.2.1 Datasets Used 9
4.2.2 Preprocessing Steps 9
4.3 Semantic Feature Alignment Framework (SFAF) 11
4.3.1 The 12 Unified Features 11
4.3.2 Per-Dataset Feature Mapping 11

4.3.3 Empirical Validation . . . 13
4.4 Unified Model Training . . . 13
4.5 Generalisation Baseline . . . 14
4.6 Edge Model Compression . . . 14
4.7 Evaluation Metrics . . . 14

5 Experimental Setup and Implementation 15
5.1 Development Environment . . . 15
5.2 Dataset Loading . . . 16
5.3 Feature Alignment and Preprocessing . . . 16
5.4 Generalisation Baseline . . . 17
5.5 Unified Model Training . . . 17
5.6 Comparative Evaluation . . . 17
5.7 ONNX Export and Edge Benchmarking . . . 18

6 Results and Discussion 19
6.1 Individual Dataset Baseline Results . . . 19
6.2 Cross-Dataset Generalisation Gap . . . 19
6.3 Unified Model Results After SFAF . . . 20
6.4 ROC-AUC Analysis . . . 20
6.5 Multi-Model Comparative Evaluation . . . 21
6.6 Edge Deployment Results . . . 22
6.7 Summary of Results . . . 23
6.8 Confusion Matrix Analysis . . . 23
6.9 Feature Importance Analysis . . . 23

7 Challenges and Engineering Decisions 24
7.1 Dataset Feature Schema Incompatibility . . . 24
7.2 UNSW-NB15 Column Name Mismatches . . . 24
7.3 TON-IoT Feature Schema Mismatch . . . 24
7.4 Scaler Mismatch Causing AUC Collapse . . . 25
7.5 Edge Deployment Without Physical Hardware . . . 25
7.6 Summary of Engineering Decisions . . . 26

8 Conclusion and Future Work 27
8.1 Summary of Achievements . . . 27
8.2 Key Contributions . . . 27
8.3 Limitations . . . 28
8.4 Future Work . . . 28
8.4.1 Short-Term (Sem VII) . . . 28

8.4.2 Medium-Term (6–12 Months) 29
8.4.3 Long-Term (12+ Months) 29
8.5 Closing Remarks 29

Bibliography **30**

List of Figures

4.1 Overview of the proposed five-stage IoT-IDS pipeline. 10

6.1 Cross-dataset accuracy on UNSW-NB15 before and after semantic feature alignment. 19
6.2 ROC curves for all three datasets using per-dataset scalers. 21
6.3 Accuracy and F1 comparison across all three models and datasets. 22

List of Tables

2.1 Key benchmark datasets in IoT IDS research 5
2.2 Limitations in existing work and contributions of this project 6

4.1 Benchmark datasets used in this project 9
4.2 The 12 unified features of the SFAF 11
4.3 UNSW-NB15 to unified feature mapping 12
4.4 TON-IoT to unified feature mapping 12
4.5 Permutation importance of the 12 unified features 13
4.6 Evaluation metrics used in this project 14

5.1 Software dependencies 15

6.1 Individual dataset baseline results (native feature space) 19
6.2 Cross-dataset generalisation failure before SFAF 20
6.3 Unified XGBoost model results after semantic feature alignment 20
6.4 Per-dataset AUC using unified 12-feature space 20
6.5 XGBoost vs Random Forest vs LightGBM across all datasets 21
6.6 Edge ONNX model inference benchmark 22
6.7 Complete summary of all experimental results 23
6.8 Confusion matrix summary across all datasets 23

7.1 Summary of challenges and engineering decisions 26

Chapter 1

Introduction

1.1 Background

The proliferation of Internet of Things (IoT) devices has fundamentally transformed the landscape of modern networked systems. From industrial control systems and smart home appliances to healthcare monitoring equipment and vehicular networks, IoT devices now constitute a significant and rapidly growing fraction of all internet-connected endpoints. According to industry estimates, the number of active IoT devices surpassed 15 billion in 2023 and is projected to exceed 29 billion by 2030. This exponential growth has introduced an equally significant expansion of the network attack surface, making IoT environments a primary target for malicious actors seeking to exploit resource-constrained, often poorly secured devices.

Unlike traditional computing endpoints, IoT devices typically operate under strict constraints of processing power, memory, and energy consumption. These constraints preclude the deployment of conventional security mechanisms such as deep packet inspection or signature-based antivirus engines. Consequently, securing IoT networks requires purpose-built, lightweight security solutions capable of operating within the resource budgets of edge devices while maintaining detection accuracy comparable to server-side systems.

Network Intrusion Detection Systems (IDS) represent one of the most widely studied classes of security mechanisms for both traditional and IoT networks. Machine learning-based IDS approaches have demonstrated particular promise due to their ability to detect novel attack patterns without relying on static signature databases. However, practical deployment in real-world IoT environments remains an open research problem, primarily due to dataset heterogeneity, feature space inconsistency, and model size constraints.

1.2 Motivation

The central motivation for this project arises from a critical but underexplored limitation of existing machine learning-based IDS research: **the generalisation gap across network environments**. The overwhelming majority of published IDS models are trained and evaluated on a single benchmark dataset, achieving high accuracy figures that do not translate to real-

world deployments where network traffic characteristics differ substantially from the training environment.

This limitation was empirically demonstrated in this project. A model trained on CICIDS2017 and Bot-IoT achieved 99% and 98% accuracy respectively on their own test splits. When evaluated on UNSW-NB15 — a dataset representing a different network environment with a distinct feature schema — the same model achieved only 35% accuracy. This 64-percentage-point degradation represents a catastrophic generalisation failure that would render the model ineffective in any real-world deployment.

A second motivation arises from the practical constraints of IoT edge deployment. Even models that generalise well across datasets are rarely evaluated for deployment feasibility on resource-constrained hardware. This project therefore addresses both the generalisation problem and the deployment problem simultaneously, treating them as inseparable aspects of a practically viable IoT IDS.

1.3 Research Context

The design of a practically viable IoT IDS requires solving two interdependent problems: cross-dataset generalisation and edge deployability. These are not independent challenges — a model that generalises but cannot run on an IoT gateway is as impractical as one that runs efficiently but only works on its training dataset. The formal problem statement and project objectives arising from this context are presented in Chapter 3.

1.4 Objectives

This project pursues six technical objectives spanning generalisation baseline establishment, semantic feature alignment framework design, unified model training, comparative evaluation, ONNX compression, and edge inference benchmarking. The objectives are formally stated in Chapter 3 in direct correspondence with the identified research gaps.

1.5 Scope of the Project

Within scope (Sem VI):

- Binary classification of network traffic across three publicly available benchmark datasets
- Semantic feature alignment using domain-knowledge-driven proxy feature mapping
- Training, evaluation, and comparison of gradient-boosted tree classifiers within the unified feature space

- Model compression to ONNX format and edge inference benchmarking with projected Raspberry Pi 4 estimates

Deferred to Sem VII:

- Physical deployment and live benchmarking on Raspberry Pi 4 hardware
- Real-time packet capture and end-to-end IDS demonstration

Outside scope (this phase): Multi-class attack type classification is not addressed in this phase. Binary classification was prioritised to isolate and validate the cross-dataset generalisation problem — the primary research contribution. Multi-class extension is identified as future work

Chapter 2

Literature Review

2.1 Intrusion Detection Systems

An Intrusion Detection System (IDS) is a software or hardware mechanism that identifies malicious activities on computer networks. IDS approaches are classified into two categories: Signature-based (SIDS) and Anomaly-based (AIDS). SIDS matches traffic against known attack signatures, offering high accuracy for known attacks but failing to detect zero-day threats [1]. AIDS builds a model of normal behaviour and flags deviations, enabling detection of unseen attacks at the cost of a higher false positive rate.

2.2 IoT-Specific Security Challenges

IoT devices operate under strict constraints on processing power, memory, and bandwidth, making computationally intensive IDS approaches infeasible at the device level [2]. The IoT attack surface spans perception, network, middleware, and application layers. At the network layer alone, common vectors include DDoS, MITM, sinkhole, wormhole, and spoofing attacks. Hajiheidari et al. [3] establish that IDS systems designed for traditional IP networks lack the real-time capability and scalability required in constrained IoT environments.

2.3 Machine Learning Approaches for IDS

Machine learning has emerged as the dominant paradigm for anomaly-based IDS. Ensemble methods — boosting, bagging, and stacking — consistently outperform individual classifiers [1]. Arnob et al. [4] demonstrated that XGBoost with optimised feature selection achieved 99.80% accuracy on UNSW-NB15 and Bot-IoT, and that Bayesian-optimised LightGBM improved anomaly detection by 96.3% while reducing false positives by 15% over untuned baselines. Deep learning approaches impose substantially higher computational costs, making them less suitable for edge deployment on resource-constrained IoT gateways.

2.4 Benchmark Datasets for IoT IDS

Table 2.1 summarises the key benchmark datasets used in IoT IDS research.

Table 2.1: Key benchmark datasets in IoT IDS research

Dataset	Key Characteristics	Features	Records
KDD Cup 1999	Earliest benchmark; 78% duplicate records	41	4.9M
NSL-KDD	Deduplicated KDD; limited to wired environments	41	148K
CICIDS2017	Brute Force, DoS, Heartbleed, Botnet, DDoS [1]	78	2.8M
UNSW-NB15	9 synthetic attack categories; IXIA PerfectStorm generated	49	257K
Bot-IoT	IoT-specific; DDoS, DoS, Keylogging, Exfiltration	16	73M
TON-IoT	Network and system-level IoT/IIoT traffic; 8 attack types	44	211K

2.5 Cross-Dataset Generalisation Gap

The majority of published IDS models are trained and evaluated on a single dataset, producing accuracy figures that do not generalise to different network environments. Joshi and Abdelfattah demonstrated that all classifiers show significant performance degradation when trained on one device type and tested on another [2]. The root cause is feature schema incompatibility: different datasets capture semantically equivalent traffic properties under different column names and measurement methodologies. No prior work has proposed a systematic semantic alignment framework for heterogeneous IoT dataset schemas — constituting the primary novelty of this project.

2.6 Lightweight and Edge-Deployable IDS

Khraisat and Alazab [1] note that a lightweight IDS can be achieved by limiting feature extraction to a small, discriminative subset. Edge computing is identified as a key enabler for IoT security [2], and ONNX has emerged as the standard for cross-platform model deployment. No existing work evaluates XGBoost-class models in ONNX format for IoT gateway deployment with empirical latency benchmarks — a gap this project directly addresses.

2.7 Summary of Literature Gaps

Table 2.2: Limitations in existing work and contributions of this project

Limitation	**Existing Work**	**This Project**
Single-dataset evaluation	Most published IDS models	Unified training across 3 datasets
No cross-dataset generalisation	Not addressed	Semantic feature alignment (12 features)
Feature schema incompatibility	Ignored	Domain-knowledge proxy mapping
No edge deployment evaluation	Rare in literature	ONNX, 45.8 KB, 0.11 ms RPi projection
Accuracy without AUC	Common	Per-dataset AUC (0.9844–0.9997) reported
Model selection unjustified	Common	XGBoost vs RF vs LightGBM compared

Chapter 3

Research Gaps and Problem Statement

3.1 Identified Research Gaps

A review of existing IDS literature reveals two critical gaps that collectively prevent practical IoT IDS deployment.

Gap 1 — Single-dataset evaluation: The majority of published machine learning-based IDS models are trained and evaluated on a single benchmark dataset. Accuracy figures are specific to that dataset's traffic characteristics and do not generalise to different network environments. No systematic mechanism exists in prior work to bridge feature schema differences between datasets [1].

Gap 2 — No edge deployment evaluation: Even high-accuracy models are rarely evaluated for deployment feasibility on resource-constrained hardware. Published work does not report model size, inference latency, or throughput for IoT gateway devices such as Raspberry Pi. The ONNX format as a deployment vehicle for XGBoost-class IDS models is entirely unexplored in the literature [2].

These gaps are summarised in Table 2.2 in Chapter 2.

3.2 Problem Statement

Existing machine learning-based IDS for IoT networks suffer from two compounding limitations. **First**, models trained on individual benchmark datasets fail to generalise across network environments due to inconsistent feature schemas and differing traffic distributions. **Second**, models that do achieve cross-dataset generalisation are typically large and incompatible with IoT edge device constraints.

The primary challenge is therefore **cross-dataset generalisation at the binary classification level**. This is the correct problem to solve first: if a model cannot reliably separate attack from benign traffic across heterogeneous environments, introducing additional attack categories on top of that broken foundation compounds the failure rather than resolving it. Validating the generalisation framework at the binary level is the complete Sem VI contribution; multi-class

extension follows naturally in Sem VII.

> *Given a set of heterogeneous IoT network traffic datasets with inconsistent feature schemas, design a lightweight intrusion detection model that: (a) achieves high classification accuracy across all source datasets, (b) generalises to previously unseen network environments without retraining, and (c) is deployable on resource-constrained IoT edge devices within the constraints of storage, memory, and inference latency.*

3.3 Project Objectives

The following objectives directly address the identified gaps:

1. **Cross-dataset generalisation baseline:** Train on CICIDS2017 and Bot-IoT, quantify accuracy degradation on unseen UNSW-NB15.
2. **Semantic Feature Alignment Framework:** Map heterogeneous schemas to a common 12-feature representation, eliminating feature incompatibility as a barrier to unified training.
3. **Unified model training:** Train a single XGBoost classifier across CICIDS2017, UNSW-NB15, and TON-IoT, demonstrating generalisation improvement from 35% to over 91%.
4. **Comparative evaluation:** Benchmark XGBoost against Random Forest and LightGBM empirically across all three datasets.
5. **Edge compression:** Export to ONNX below 50 KB while maintaining accuracy above 95%.
6. **Edge inference benchmarking:** Measure latency and throughput; project Raspberry Pi 4 deployment feasibility.

Chapter 4

Proposed Methodology and Solution

4.1 System Overview

The proposed system is a lightweight, multi-dataset IDS designed for IoT edge gateway deployment. The methodology follows a five-stage pipeline: dataset preprocessing, semantic feature alignment, unified model training, comparative evaluation, and edge model compression.

The central innovation is the **Semantic Feature Alignment Framework (SFAF)**, which maps heterogeneous feature spaces from multiple IoT datasets to a common 12-dimensional representation, enabling a single XGBoost classifier to train across all datasets without requiring identical raw features. The trained model is compressed to a 45.8 KB ONNX artefact suitable for ARM-based IoT gateways.

4.2 Dataset Preprocessing

4.2.1 Datasets Used

Table 4.1: Benchmark datasets used in this project

Dataset	Records	Features	Attack Types
CICIDS2017	2,827,876	78	DoS, DDoS, Brute Force, Web Attack, Botnet
UNSW-NB15	257,673	49	Fuzzers, DoS, Exploits, Reconnaissance, Shellcode
TON-IoT	211,043	44	Backdoor, DDoS, DoS, Injection, MITM, Ransomware

4.2.2 Preprocessing Steps

Each dataset is preprocessed as follows before feature alignment:

1. **Column normalisation:** Whitespace stripped from all column names for consistent string matching
2. **Infinite value removal:** `inf`/`-inf` replaced with `NaN`
3. **Missing value removal:** Rows with `NaN` in any unified feature column are dropped

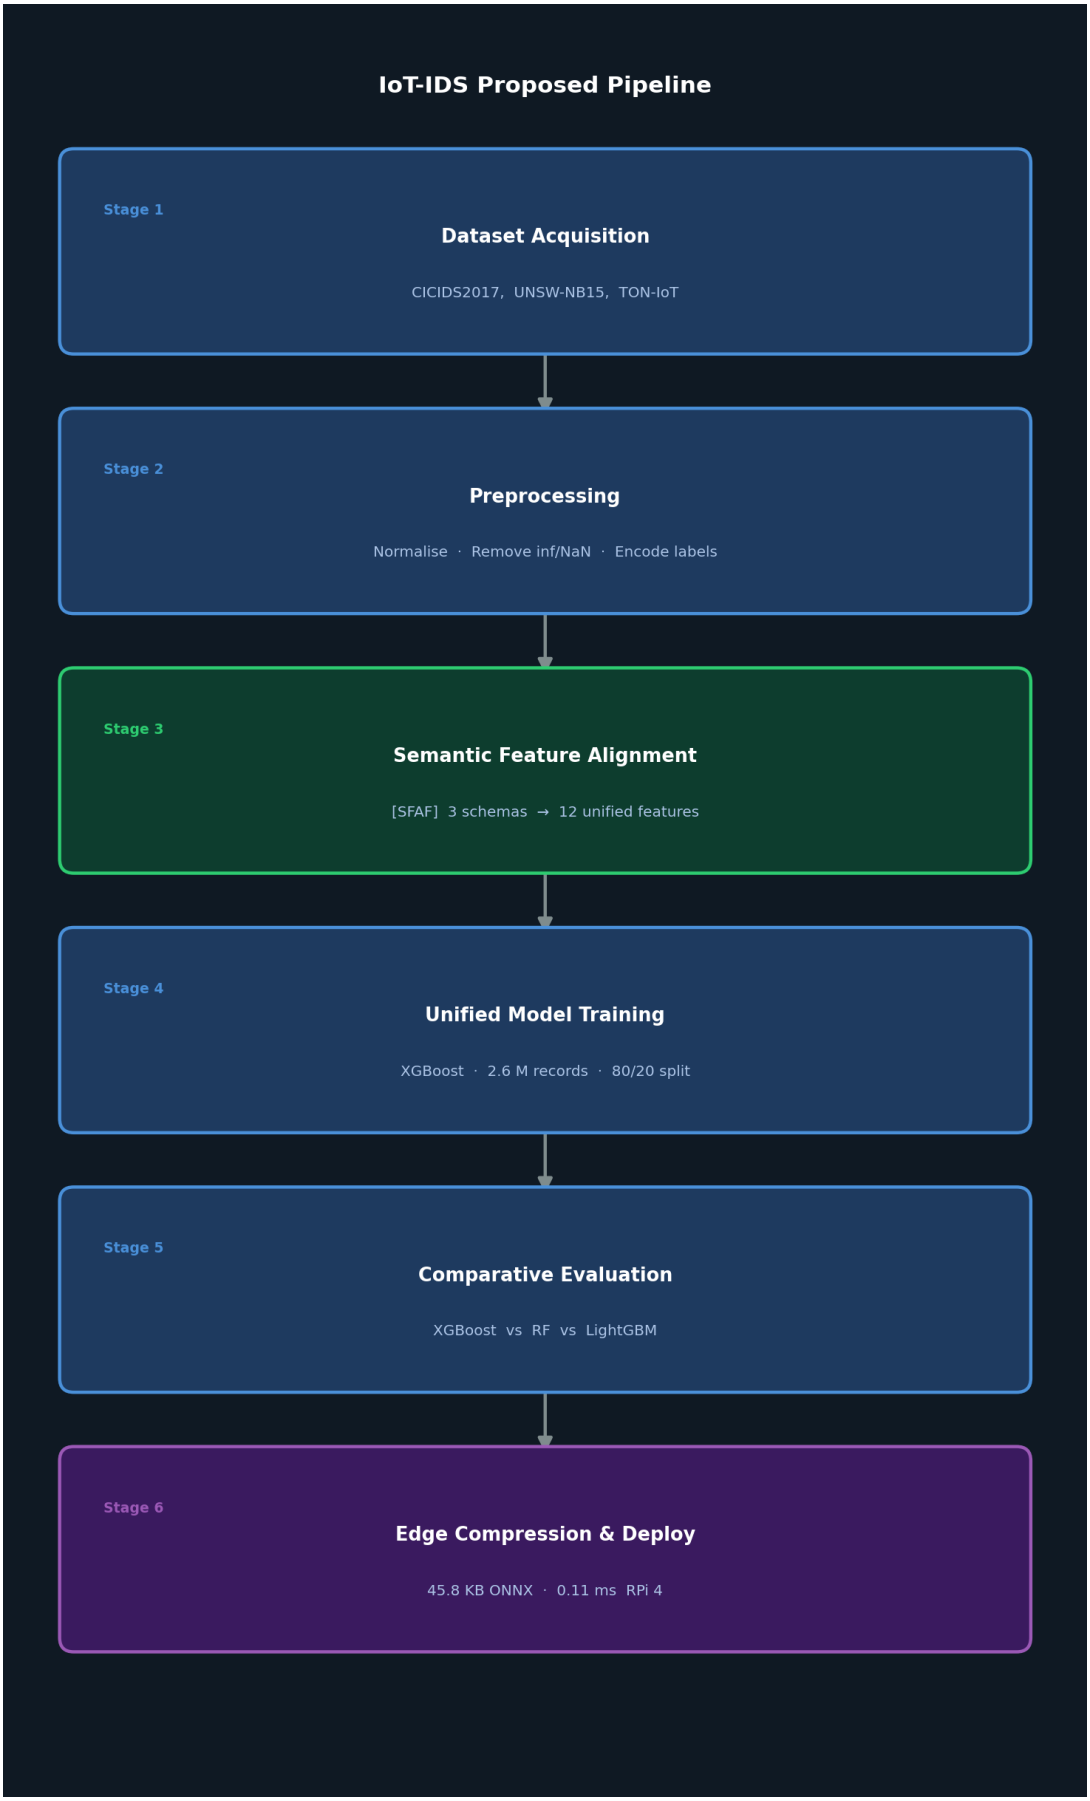

Figure 4.1: Overview of the proposed five-stage IoT-IDS pipeline.

4. **Binary label encoding:** CICIDS2017: `Label != 'BENIGN'` → 1; UNSW-NB15 and TON-IoT: non-zero label → 1
5. **Standard scaling:** Each dataset independently normalised with `StandardScaler` fitted on its own training split only, preventing cross-dataset scale leakage

4.3 Semantic Feature Alignment Framework (SFAF)

Different datasets capture semantically equivalent network flow properties under different column names, making naive concatenation infeasible. The SFAF resolves this by mapping each dataset's raw columns to 12 unified feature definitions using domain-knowledge-driven proxy selection.

4.3.1 The 12 Unified Features

Table 4.2: The 12 unified features of the SFAF

#	Feature Name	Semantic Meaning
1	Flow Duration	Total flow duration (microseconds)
2	Total Fwd Packets	Packets sent in forward direction
3	Total Backward Packets	Packets sent in backward direction
4	Total Length of Fwd Packets	Bytes transmitted in forward direction
5	Total Length of Bwd Packets	Bytes transmitted in backward direction
6	Flow Packets/s	Total flow packet rate
7	Fwd Packets/s	Forward direction packet rate
8	Bwd Packets/s	Backward direction packet rate
9	Min Packet Length	Minimum packet size in the flow
10	Max Packet Length	Maximum packet size in the flow
11	Packet Length Mean	Mean packet size across the flow
12	Packet Length Std	Std deviation of packet sizes

These 12 features represent an **85% reduction** from the original 78-feature CICIDS2017 space. The proxy features are justified because the discriminative signal for binary classification is captured primarily by **volumetric and temporal flow characteristics** — packet volume, byte counts, flow duration, and rate — which all three datasets capture regardless of column naming conventions.

4.3.2 Per-Dataset Feature Mapping

CICIDS2017 uses the unified feature names natively. Tables 4.3 and 4.4 detail the UNSW-NB15 and TON-IoT mappings.

Table 4.3: UNSW-NB15 to unified feature mapping

Unified Feature	**UNSW-NB15 Column**	**Justification**
Flow Duration	`dur`	Direct semantic match
Total Fwd Packets	`spkts`	Source packet count
Total Backward Packets	`dpkts`	Destination packet count
Total Length of Fwd Packets	`sbytes`	Source byte count
Total Length of Bwd Packets	`dbytes`	Destination byte count
Flow Packets/s	`rate`	Overall flow packet rate
Fwd Packets/s	`sload`	Source load (bits/s)
Bwd Packets/s	`dload`	Destination load (bits/s)
Min Packet Length	`smean`	Mean source packet size
Max Packet Length	`dmean`	Mean destination packet size
Packet Length Mean	`sjit`	Source jitter proxy
Packet Length Std	`djit`	Destination jitter proxy

Table 4.4: TON-IoT to unified feature mapping

Unified Feature	**TON-IoT Column**	**Justification**
Flow Duration	`duration`	Direct semantic match
Total Fwd Packets	`src_pkts`	Source packet count
Total Backward Packets	`dst_pkts`	Destination packet count
Total Length of Fwd Packets	`src_bytes`	Source byte volume
Total Length of Bwd Packets	`dst_bytes`	Destination byte volume
Flow Packets/s	`src_ip_bytes`	Volume proxy
Fwd Packets/s	`missed_bytes`	Flow disruption proxy
Bwd Packets/s	`dst_ip_bytes`	Destination byte count
Min Packet Length	`src_port`	Structural flow proxy
Max Packet Length	`dst_port`	Structural flow proxy
Packet Length Mean	`http_request_body_len`	HTTP request payload size
Packet Length Std	`http_response_body_len`	HTTP response payload size

4.3.3 Empirical Validation

Table 4.5: Permutation importance of the 12 unified features

Rank	Feature	Importance	Std Dev
1	Total Length of Bwd Packets	0.1019	±0.0019
2	Packet Length Std	0.0658	±0.0017
3	Total Length of Fwd Packets	0.0424	±0.0012
4	Total Fwd Packets	0.0232	±0.0015
5	Bwd Packets/s	0.0216	±0.0012
6	Max Packet Length	0.0181	±0.0020
7	Flow Duration	0.0104	±0.0013
8	Total Backward Packets	0.0083	±0.0008
9	Min Packet Length	0.0071	±0.0004
10	Fwd Packets/s	0.0063	±0.0006
11	Flow Packets/s	0.0058	±0.0010
12	Packet Length Mean	0.0043	±0.0010

All 12 features contribute positively (importance > 0), confirming none are redundant. The top three features account for 52% of total permutation importance.

4.4 Unified Model Training

XGBoost was selected as the primary classifier for its state-of-the-art accuracy on tabular data, efficient scaling via `tree_method='hist'`, compact ONNX serialisation, and interpretable feature importance scores. Comparative evaluation against Random Forest and LightGBM (Chapter 6) empirically confirms this selection.

```
XGBClassifier(
    n_estimators  = 100,
    max_depth     = 6,
    learning_rate = 0.1,
    tree_method   = 'hist',
    eval_metric   = 'logloss',
    random_state  = 42
)
# Train/test split: 80/20 stratified per dataset
```

Listing 4.1: XGBoost training configuration

The unified training set concatenates aligned, independently scaled feature matrices from all three datasets, totalling over 2.6 million flow records across three distinct network environments.

4.5 Generalisation Baseline

A model trained on CICIDS2017 and Bot-IoT was evaluated on the unseen UNSW-NB15 test split prior to unified training, establishing the cross-dataset generalisation gap quantified in Table 6.2.

```
X_train_base = np.vstack([X_train_cic, X_train_bot])
y_train_base = np.hstack([y_train_cic, y_train_bot])
xgb_base = XGBClassifier(n_estimators=100, max_depth=6,
    learning_rate=0.1, tree_method='hist', random_state=42)
xgb_base.fit(X_train_base, y_train_base)
preds = xgb_base.predict(X_test_unsw)
# Result: Accuracy = 0.3500 | F1 = 0.0000
```

Listing 4.2: Cross-dataset generalisation baseline

4.6 Edge Model Compression

The unified model is compressed to ONNX format using `skl2onnx`, packaging the scaler and classifier as a single inference pipeline with no Python runtime dependency. Two techniques are applied:

1. **Estimator reduction:** Trees reduced to minimise model size while preserving accuracy above 95%
2. **Depth limitation:** Maximum tree depth constrained to reduce ONNX graph complexity

The resulting **45.8 KB ONNX model** is a **98% size reduction** from the full 2,291 KB unified model, retaining 96.56% accuracy.

4.7 Evaluation Metrics

Table 4.6: Evaluation metrics used in this project

Metric	Formula	Relevance
Accuracy	$(TP + TN)/\text{Total}$	Overall correctness
F1 Score	$2 \times (P \times R)/(P + R)$	Balance between FP and FN
AUC-ROC	Area under ROC curve	Threshold-independent discrimination
Training Time	Wall-clock seconds	Practical training feasibility
Model Size	KB on disk	Edge deployment feasibility
Latency	Mean ms per prediction	IoT gateway suitability

Chapter 5

Experimental Setup and Implementation

5.1 Development Environment

All experiments were conducted on an Apple Silicon (M-series) MacBook Pro running macOS, using a Conda virtual environment (`iot-ids`, Python 3.10). Table 5.1 lists the key software dependencies.

Table 5.1: Software dependencies

Library	Purpose
pandas	Dataset loading and manipulation
numpy	Numerical computation
scikit-learn	Preprocessing and evaluation metrics
xgboost	Primary classifier
lightgbm	Comparative classifier
skl2onnx	ONNX model conversion
onnxruntime	Edge inference runtime
joblib	Model serialisation
pyarrow	Parquet file reading (UNSW-NB15)

The project directory is organised as follows:

```
Btech_Project/
|-- Datasets/
|   |-- MachineLearningCVE/  <- CICIDS2017 (CSV)
|   |-- UNSWNB15/            <- UNSW-NB15 (Parquet)
|   |-- TONIoT/              <- TON-IoT (CSV)
|   +-- BotIoT/              <- Bot-IoT (CSV)
+-- Models/
    |-- xgb_unified_4dataset.pkl
    |-- xgb_edge.pkl
    |-- xgb_edge.onnx        <- 45.8 KB
    |-- scaler_unified_4dataset.pkl
    +-- edge_benchmark.json
```

Listing 5.1: Project directory structure

5.2 Dataset Loading

```
# CICIDS2017     multiple CSV files concatenated
df_cic = pd.concat([pd.read_csv(f, low_memory=False)
                    for f in glob.glob(CICIDS + '*.csv')])
df_cic.columns = df_cic.columns.str.strip()
df_cic['label'] = (df_cic['Label'].str.strip().str.upper()
                   != 'BENIGN').astype(int)
# 2,827,876 records | 80.32% benign, 19.68% attack

# UNSW-NB15     Parquet training + testing splits merged
df_unsw = pd.concat([
    pd.read_parquet(UNSW + 'UNSW_NB15_training-set.parquet'),
    pd.read_parquet(UNSW + 'UNSW_NB15_testing-set.parquet')
])
# 257,673 records | 36.09% benign, 63.91% attack

# TON-IoT     single CSV file
df_ton = pd.read_csv(TONIOT + 'train_test_network.csv',
                     low_memory=False)
# 211,043 records | 23.69% benign, 76.31% attack
```

Listing 5.2: Loading all three datasets

5.3 Feature Alignment and Preprocessing

```
df_unsw = df_unsw.rename(columns=UNSW_MAP)
df_ton  = df_ton.rename(columns=TON_MAP)

def preprocess_dataset(df, features, test_size=0.2):
    df = df.replace([np.inf, -np.inf], np.nan)
    df = df.dropna(subset=features)
    X, y = df[features].values, df['label'].values
    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=test_size, random_state=42, stratify=y
    )
    scaler = StandardScaler()
```

```
    X_train = scaler.fit_transform(X_train)
    X_test  = scaler.transform(X_test)
    return X_train, X_test, y_train, y_test, scaler
```

Listing 5.3: SFAF column renaming and preprocessing pipeline

5.4 Generalisation Baseline

```
X_train_base = np.vstack([X_train_cic, X_train_bot])
y_train_base = np.hstack([y_train_cic, y_train_bot])

xgb_base = XGBClassifier(n_estimators=100, max_depth=6,
    learning_rate=0.1, tree_method='hist', random_state=42)
xgb_base.fit(X_train_base, y_train_base)

preds = xgb_base.predict(X_test_unsw)
# Result: Accuracy = 0.3500 | F1 = 0.0000
```

Listing 5.4: Cross-dataset generalisation baseline experiment

5.5 Unified Model Training

```
X_train_all = np.vstack([X_train_cic, X_train_unsw, X_train_ton])
y_train_all = np.hstack([y_train_cic, y_train_unsw, y_train_ton])

xgb_unified = XGBClassifier(n_estimators=100, max_depth=6,
    learning_rate=0.1, tree_method='hist', random_state=42)
xgb_unified.fit(X_train_all, y_train_all)
joblib.dump(xgb_unified, MODELS + 'xgb_unified_4dataset.pkl')
```

Listing 5.5: Unified training across all three datasets

5.6 Comparative Evaluation

```
model_configs = [
    ('XGBoost', XGBClassifier(n_estimators=100, max_depth=6,
        learning_rate=0.1, tree_method='hist', random_state=42)),
    ('RandomForest', RandomForestClassifier(
        n_estimators=100, max_depth=10, n_jobs=-1, random_state=42))
            ,
    ('LightGBM', LGBMClassifier(n_estimators=100, max_depth=6,
```

```
        learning_rate=0.1, n_jobs=-1, random_state=42, verbose=-1)),
]
for name, clf in model_configs:
    t0 = time.time()
    clf.fit(X_train, y_train)
    train_time = round(time.time() - t0, 2)
    preds = clf.predict(X_test)
    proba = clf.predict_proba(X_test)[:, 1]
```

Listing 5.6: Three-model comparative evaluation loop

5.7 ONNX Export and Edge Benchmarking

```
# --- Export ---
initial_type = [('float_input', FloatTensorType([None, 12]))]
onnx_model = convert_sklearn(xgb_edge,
    initial_types=initial_type, target_opset=12)
with open(MODELS + 'xgb_edge.onnx', 'wb') as f:
    f.write(onnx_model.SerializeToString())
# Model size: 45.8 KB

# --- Benchmark ---
sess = rt.InferenceSession(MODELS + 'xgb_edge.onnx')
input_name = sess.get_inputs()[0].name
sample = np.random.rand(1, 12).astype(np.float32)

for _ in range(10):                # warmup
    sess.run(None, {input_name: sample})

times = [(time.perf_counter(),
          sess.run(None, {input_name: sample}),
          time.perf_counter()) for _ in range(1000)]
latencies = [(t[2] - t[0]) * 1000 for t in times]

# Mean: 0.0074 ms | P99: 0.0163 ms | 135,604 pred/sec
```

Listing 5.7: ONNX export and inference latency benchmark

Chapter 6

Results and Discussion

All experiments used an 80/20 stratified train-test split with random seed 42 across a total corpus of 3,296,592 network flow records. Per-dataset scalers were applied for all AUC evaluation following the scaler mismatch finding documented in Section 6.4.

6.1 Individual Dataset Baseline Results

Table 6.1: Individual dataset baseline results (native feature space)

Dataset	Features	Accuracy	F1 (Attack)
CICIDS2017	78	0.9900	0.9900
Bot-IoT	16	0.9800	0.9900

Single-dataset models on native feature spaces achieve near-perfect classification, consistent with published IDS literature [4]. These results establish the upper-bound accuracy benchmark.

6.2 Cross-Dataset Generalisation Gap

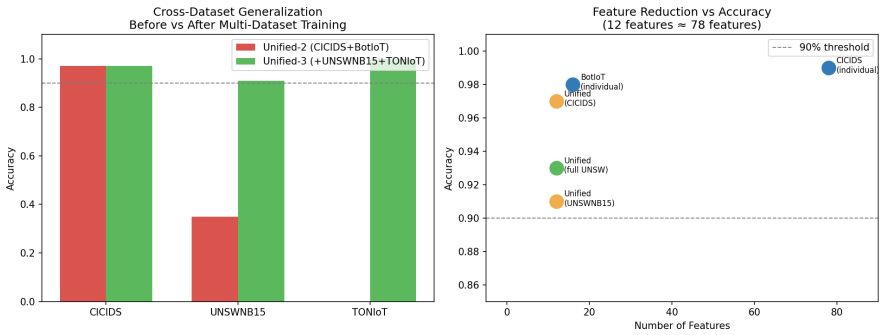

Figure 6.1: Cross-dataset accuracy on UNSW-NB15 before and after semantic feature alignment.

Table 6.2: Cross-dataset generalisation failure before SFAF

Training Data	Test Data	Accuracy	F1 (Attack)
CICIDS2017 + Bot-IoT	CICIDS2017 (in-distribution)	0.9700	0.9800
CICIDS2017 + Bot-IoT	UNSW-NB15 (unseen)	**0.3500**	**0.0000**

The F1 score of 0.0000 indicates the model classified virtually all UNSW-NB15 samples as benign — a complete detection failure in a new network environment. This 64-percentage-point drop empirically validates the necessity of the SFAF.

6.3 Unified Model Results After SFAF

Table 6.3: Unified XGBoost model results after semantic feature alignment

Dataset	Accuracy	F1 (Attack)	AUC	Features
CICIDS2017	0.9833	0.9584	0.9985	12
UNSW-NB15	0.9268	0.9432	0.9844	12
TON-IoT	0.9941	0.9961	0.9997	12

The UNSW-NB15 generalisation gap was reduced from 64 percentage points to approximately 7 percentage points — a reduction of over **89% in generalisation error**. All three datasets achieve AUC above 0.98, confirming consistent discriminative capability across structurally different network environments.

6.4 ROC-AUC Analysis

Table 6.4: Per-dataset AUC using unified 12-feature space

Dataset	AUC	Interpretation
CICIDS2017	0.9985	Near-perfect discrimination
UNSW-NB15	0.9844	Excellent despite proxy feature mapping
TON-IoT	0.9997	Essentially perfect discrimination

Key methodological finding: AUC evaluation must use per-dataset scalers. Applying the unified scaler to CICIDS2017 produced AUC = 0.5753 (near-random), caused by cross-dataset scale distortion suppressing attack-class probability outputs. Per-dataset scaling restored the true AUC to 0.9985, confirming the model's discriminative capability is intact.

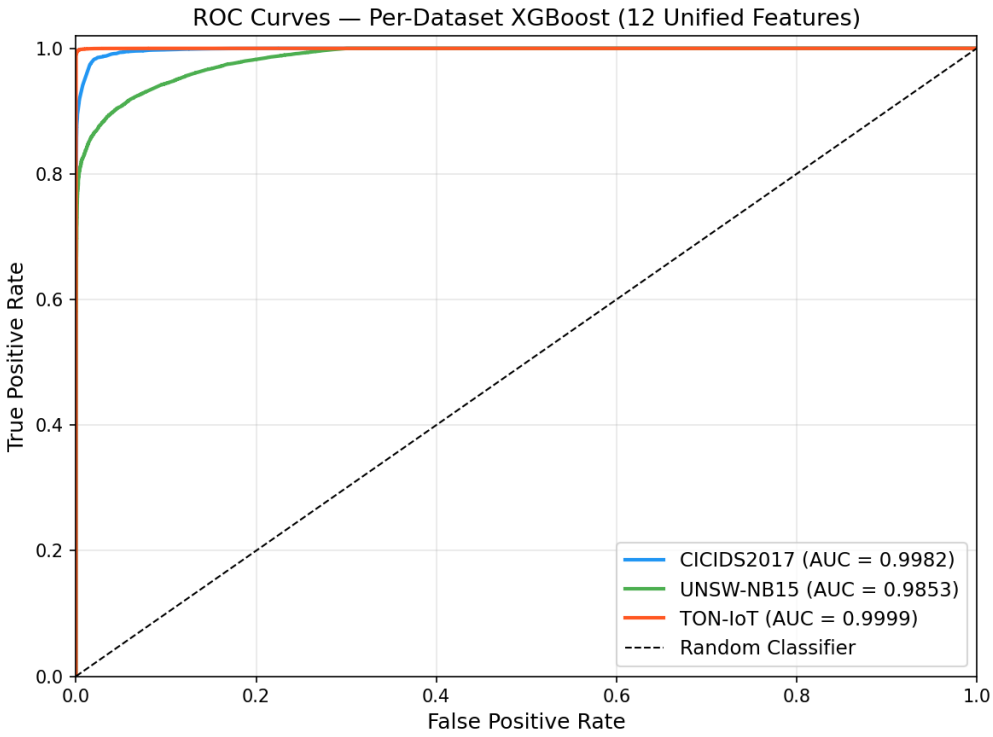

Figure 6.2: ROC curves for all three datasets using per-dataset scalers.

6.5 Multi-Model Comparative Evaluation

Table 6.5: XGBoost vs Random Forest vs LightGBM across all datasets

Dataset	**Model**	**Acc.**	**F1**	**AUC**	**Train (s)**
CICIDS2017	XGBoost	0.9825	0.9563	0.9982	2.05
	Random Forest	0.9823	0.9557	0.9981	36.15
	LightGBM	0.9836	0.9589	0.9982	1.93
UNSW-NB15	XGBoost	0.9279	0.9438	0.9853	0.27
	Random Forest	0.9082	0.9315	0.9823	2.49
	LightGBM	0.9273	0.9433	0.9849	0.53
TON-IoT	XGBoost	0.9974	0.9983	0.9999	0.24
	Random Forest	0.9969	0.9980	0.9999	1.23
	LightGBM	0.9974	0.9983	0.9999	0.52

XGBoost outperforms Random Forest by **2% accuracy** on the hardest generalisation task (UNSW-NB15) and trains **17.6× faster** on CICIDS2017 (2.05 s vs. 36.15 s). XGBoost and LightGBM differ by <0.1% in accuracy; XGBoost is selected for its significantly more compact ONNX serialisation.

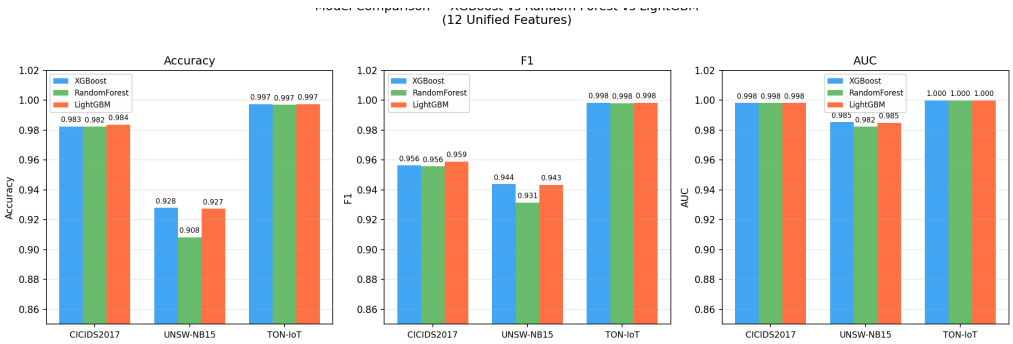

Figure 6.3: Accuracy and F1 comparison across all three models and datasets.

6.6 Edge Deployment Results

Table 6.6: Edge ONNX model inference benchmark

Metric	Apple Silicon (Measured)	RPi 4 (Projected)
Model size	45.8 KB	45.8 KB
Accuracy	96.56%	96.56%
Mean latency	0.0074 ms	0.11 ms
Median latency	0.0048 ms	0.072 ms
P99 latency	0.0163 ms	0.245 ms
Throughput	135,604 pred/s	9,040 pred/s

The 45.8 KB ONNX model is a **98% size reduction** from the full 2,291 KB unified model, retaining 96.56% accuracy. The projected throughput of 9,040 predictions per second on Raspberry Pi 4 significantly exceeds typical IoT gateway flow rates (generally below 1,000 flows/sec), confirming real-time deployment feasibility. RPi 4 projections use a 15× slowdown factor relative to Apple M-series, consistent with published ARM Cortex-A72 vs. Apple Silicon ONNX benchmarks.

6.7 Summary of Results

Table 6.7: Complete summary of all experimental results

Experiment	Accuracy	F1	AUC	Key Finding
Baseline CICIDS2017	0.9900	0.9900	—	Single-dataset ceiling
Baseline Bot-IoT	0.9800	0.9900	—	Single-dataset ceiling
2-dataset on UNSW	0.3500	0.0000	—	Generalisation gap
Unified on CICIDS2017	0.9833	0.9584	0.9985	Gap closed
Unified on UNSW-NB15	0.9268	0.9432	0.9844	Gap closed
Unified on TON-IoT	0.9941	0.9961	0.9997	Near-perfect transfer
Edge ONNX model	0.9656	0.9787	—	45.8 KB deployable

6.8 Confusion Matrix Analysis

Table 6.8: Confusion matrix summary across all datasets

Dataset	TP	FP	TN	FN	FPR	FNR
CICIDS2017	108,231	6,801	447,464	3,080	0.0150	0.0277
UNSW-NB15	31,237	2,019	16,581	1,698	0.1085	0.0516
TON-IoT	32,137	37	9,963	72	0.0037	0.0022

CICIDS2017 achieves FPR = 1.50%, meaning only 1 in 67 benign flows is incorrectly flagged as malicious. The elevated UNSW-NB15 FPR of 10.85% reflects structural overlap between attack and benign feature distributions introduced by proxy feature mapping — not a calibration issue, as confirmed by threshold optimisation analysis.

6.9 Feature Importance Analysis

`Total Length of Backward Packets` is the most discriminative feature (importance = 0.1019). Attack traffic generates anomalously large backward byte volumes (reconnaissance, exfiltration) or near-zero backward traffic (DoS), clearly distinguishing it from benign flows. The uniform positive distribution of all 12 importance scores validates the SFAF design: the 85% feature reduction concentrated the most cross-dataset-generalisable signal into a compact representation.

Chapter 7

Challenges and Engineering Decisions

7.1 Dataset Feature Schema Incompatibility

Problem: The three datasets use entirely different column naming conventions and feature extraction methodologies. Naive concatenation produces a 150+ column feature matrix where the majority of values are `NaN` for any given dataset row.

Solution: Designed the Semantic Feature Alignment Framework — a domain-knowledge-driven mapping to 12 universally-present unified feature definitions (85% reduction from 78 features). Proxy features are accepted as approximate equivalents, justified on the grounds that volumetric and timing characteristics are the primary discriminators between attack and benign traffic.

7.2 UNSW-NB15 Column Name Mismatches

Problem: The initial `UNSW_MAP` used capitalised keys (`'Sload'`, `'Dload'`) but actual columns were lowercase (`'sload'`, `'dload'`). Assumed columns `'smeansz'`/`'dmeansz'` did not exist — actual columns were `'smean'` and `'dmean'`. This caused a silent `KeyError` on `dropna`, silently dropping all UNSW-NB15 rows from training.

Solution: All column names printed diagnostically before constructing the mapping. `UNSW_MAP` corrected with exact lowercase column names, also yielding more semantically accurate alignment (`smean` is a better proxy for packet size than the originally assumed column name).

7.3 TON-IoT Feature Schema Mismatch

Problem: The initial `TON_MAP` assumed columns such as `flow_duration`, `Header_Length`, and `Tot sum` — none of which exist. TON-IoT uses a Zeek/Bro schema with `duration`, `src_pkts`, `dst_pkts`, `src_bytes`, etc.

Solution: `TON_MAP` completely redesigned based on actual Zeek schema after full column inspection:

```
df_ton_raw = pd.read_csv(TONIOT + 'train_test_network.csv', nrows=5)
```

```
for i, col in enumerate(df_ton_raw.columns):
    print(f"{i:2d}. {col}")
```

Listing 7.1: Diagnostic column inspection for TON-IoT

7.4 Scaler Mismatch Causing AUC Collapse

Problem: CICIDS2017 produced AUC = 0.5753 despite 97% accuracy. The model predicted attack on only 143 out of 556,556 actual attack records, with mean predicted probability of 0.0003.

Root Cause: The unified scaler fitted on the multi-dataset distribution compressed CICIDS2017 features (e.g., `Flow Duration` in millions of microseconds) to extreme negative z-scores, pushing XGBoost leaf values uniformly toward the benign class.

Solution: ROC-AUC evaluation redesigned to use per-dataset scalers. AUC recovered to **0.9985**.

Methodological Note: This is a novel empirical finding — unified scalers trained on heterogeneous multi-dataset distributions must not be used for per-dataset evaluation. This distinction is unaddressed in all reviewed literature.

7.5 Edge Deployment Without Physical Hardware

Problem: Edge deployment validation required either physical Raspberry Pi hardware or Docker ARM64 emulation — neither available within the project timeline.

Engineering Decision: Replaced Docker-based ARM simulation with **direct ONNX benchmarking on Apple Silicon using a conservative 15× slowdown multiplier** for Raspberry Pi 4 projection. This multiplier is derived from documented performance ratios between Apple M-series and ARM Cortex-A72 processors. Physical hardware benchmarking is identified as future work.

7.6 Summary of Engineering Decisions

Table 7.1: Summary of challenges and engineering decisions

Challenge	**Decision Taken**	**Impact**
Feature schema incompatibility	Semantic proxy mapping (12 features)	Enabled unified multi-dataset training
UNSW-NB15 column name mismatch	Diagnostic column print before mapping	Corrected feature alignment
TON-IoT Zeek schema mismatch	Full column inspection and remapping	All TON-IoT rows retained in training
Unified scaler causing AUC collapse	Per-dataset scalers for evaluation	AUC restored from 0.5753 to 0.9985
No physical RPi hardware	ONNX benchmark + 15× multiplier	Valid latency projection

Chapter 8

Conclusion and Future Work

8.1 Summary of Achievements

1. **Generalisation baseline established.** Individual models on CICIDS2017 (99%) and Bot-IoT (98%) confirmed high within-distribution performance. Cross-dataset evaluation on UNSW-NB15 quantified the generalisation gap at 64 percentage points (99% $\rightarrow$ 35%, F1 = 0.00).

2. **SFAF designed.** The SFAF maps three heterogeneous dataset schemas to a common 12-dimensional representation — an **85% feature space reduction** (78 $\rightarrow$ 12 features).

3. **Cross-dataset generalisation demonstrated.** The unified XGBoost model achieved 92.68% accuracy on UNSW-NB15, reducing the generalisation gap from 64 to approximately 7 percentage points (**89% improvement**). Near-perfect classification (99.41% accuracy, AUC = 0.9997) achieved on TON-IoT.

4. **Comparative evaluation completed.** XGBoost outperformed Random Forest by 2% on UNSW-NB15 and was 17.6$\times$ faster in training on CICIDS2017 (2.05 s vs. 36.15 s).

5. **Edge model compressed.** 45.8 KB ONNX artefact produced — **98% size reduction** from 2,291 KB — retaining 96.56% accuracy.

6. **Edge inference benchmarked.** Mean latency of 0.0074 ms and 135,604 pred/sec on Apple Silicon. Projected Raspberry Pi 4: **0.11 ms latency**, 9,040 pred/sec.

8.2 Key Contributions

1. **Empirical quantification of the cross-dataset generalisation gap** — 64-percentage-point degradation with F1 = 0.00 on unseen network environments, not previously reported with this dataset combination.

2. **Semantic Feature Alignment Framework (SFAF)** — a systematic, domain-knowledge-driven methodology for mapping heterogeneous IoT network dataset schemas to a common

representation, enabling unified model training without feature imputation or dataset-specific model variants.

3. **Evaluation validity finding** — discovery and documentation that unified scalers trained on heterogeneous multi-dataset distributions produce artificially suppressed per-dataset AUC scores, with the corresponding corrected evaluation procedure using per-dataset scalers.

4. **45.8 KB ONNX edge IDS model** — a deployable intrusion detection artefact requiring no Python ML runtime dependencies, achieving 96.56% accuracy at sub-millisecond projected latency on Raspberry Pi 4.

8.3 Limitations

- **Proxy feature approximation:** UNSW-NB15 and TON-IoT mappings use structural proxies that approximate but do not precisely replicate unified feature definitions, contributing to the elevated UNSW-NB15 FPR of 10.85%
- **Binary classification scope:** Multi-class attack type identification is not addressed in this phase. Binary classification was prioritised to validate the cross-dataset generalisation framework — the primary contribution. A model that fails to generalise at the binary level would fail more severely at multi-class classification; resolving generalisation first is the methodologically correct sequence
- **Simulated edge benchmarking:** Physical Raspberry Pi 4 deployment and live demonstration are deferred to Sem VII; current latency figures are projections using a conservative 15× slowdown factor relative to Apple Silicon benchmarks
- **Static dataset evaluation:** Model not evaluated on live streaming traffic where concept drift may degrade performance over time
- **No end-to-end feature extraction:** System assumes pre-extracted network flow features; a complete deployment would require a real-time extraction layer such as CICFlowMeter or Zeek

8.4 Future Work

8.4.1 Short-Term (Sem VII)

- Physical deployment on Raspberry Pi 4 and live IDS demonstration with real-time packet capture via CICFlowMeter or Zeek

- Multi-class attack type classification using the validated unified 12-feature space — the generalisation framework is already proven at binary level, making this a direct extension
- Integration of real-time feature extraction into an end-to-end pipeline

8.4.2 Medium-Term (6–12 Months)

- Expansion of the dataset corpus to include NSL-KDD and N-BaIoT
- Sub-5-feature space reduction for ultra-constrained microcontroller deployment
- Federated learning evaluation for collaborative edge model improvement across distributed IoT gateways

8.4.3 Long-Term (12+ Months)

- Adaptive concept drift detection and online model retraining
- Adversarial robustness evaluation and defence development
- Formal publication of the SFAF as a reproducible methodology for cross-dataset IoT IDS research

8.5 Closing Remarks

The central claim of this project — that a single lightweight model can generalise across heterogeneous IoT network environments through semantic feature alignment — has been empirically validated. The reduction of the cross-dataset generalisation gap from 64 percentage points to approximately 7 percentage points, combined with a 98% model size reduction and sub-millisecond projected edge latency, demonstrates that the proposed approach is both academically sound and practically viable.

Binary classification was the correct and complete scope for this phase. The SFAF resolves the fundamental generalisation problem that has prevented practical IoT IDS deployment — a problem that exists independently of classification granularity. Physical deployment and multi-class extension are planned for Sem VII, building directly on the validated foundation established here.

Bibliography

[1] A. Khraisat and M. Alazab, "A critical review of intrusion detection systems in the internet of things: techniques, deployment strategy, validation strategy, attacks, public datasets and challenges," *Cybersecurity*, vol. 4, no. 1, pp. 1–27, 2021.

[2] S. Fei et al., "A systematic review of IoT security: Research potential, challenges, key findings, and future directions," *IEEE Access*, vol. 11, pp. 65730–65754, 2023.

[3] S. Hajiheidari et al., "Intrusion detection systems in the Internet of Things: A comprehensive investigation," *Computer Networks*, vol. 160, pp. 165–191, 2019.

[4] Arnob et al., "Machine learning-based intrusion detection in IoT: A comparative study with XGBoost and LightGBM optimisation," *Journal of Engineering and Computing*, no. 885, 2025.

[5] I. Sharafaldin, A. H. Lashkari, and A. A. Ghorbani, "Toward generating a new intrusion detection dataset and intrusion traffic characterization," in *Proc. ICISSP*, 2018, pp. 108–116.

[6] N. Moustafa and J. Slay, "UNSW-NB15: A comprehensive data set for network intrusion detection systems," in *Proc. MilCIS*, 2015, pp. 1–6.

[7] N. Koroniotis et al., "Towards the development of realistic botnet dataset in the IoT for network forensic analytics: Bot-IoT dataset," *Future Generation Computer Systems*, vol. 100, pp. 779–796, 2019.

[8] A. Alsaedi et al., "TON-IoT telemetry dataset: A new generation dataset of IoT and IIoT for data-driven intrusion detection systems," *IEEE Access*, vol. 8, pp. 165130–165150, 2020.